


Real-data fit and conditional lifecycle — no hardware price or task-quality claim

Finite-grid fit; holdout RMSE 0.01934



Target loss 2.9; training + inference proxy

