

One successful request: input → model execution → output

A Task B Tokenize / submit C Organize D Operators E Resources F Decode / check

1 Application / task

A, F: define task and validate answer

2 Model / workload

A, B, F: messages, token IDs, decoded text

3 Inference system

B, C: accept and organize execution

4 Operators / runtime

D: execute model operations

5 Processor / storage

E: CPU, GPU, device memory, saved records

6 Interconnect / data center

No cross-host measurement supplied

Roles can span layers. This diagram does not show per-stage timing or physical traffic.